

CFRM 505 Homework 3

Eunki Chung
eunkich@uw.edu

February 3, 2023

Problem 1

Implement the following three methods to generate 100,000 i.i.d. samples of $X \sim N(0, 1)$. In each case, plot a histogram of your samples and estimate the average number of uniform random variables required per sample and the average computation time required per sample.

- 1) Acceptance-rejection with $g(x) = 0.5e^x \mathbf{1}_{(-\infty, 0]}(x) + 0.5e^{-x} \mathbf{1}_{(0, \infty)}(x)$ and $a = \sqrt{2e/\pi}$.
- 2) The Box-Muller algorithm. Be sure to use all of the samples you generate.
- 3) The Marsaglia-Bray algorithm. Be sure to use all of the samples you generate.

You can time a block of python code with the function `time.perf_counter`. For example,

1 Acceptance-rejection

```
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt

np.random.seed(12)
```

$$g(x) = 0.5e^x \mathbf{1}_{x < 0} + 0.5e^{-x} \mathbf{1}_{x \geq 0}$$

Recall the following property of an Exponential r.v. X , if

$$X \sim \text{Exp}(\lambda)$$

$$f_X(x) = \begin{cases} \lambda e^{-\lambda x} & x \geq 0 \\ 0 & o.w. \end{cases}$$

$$X = -\frac{1}{\lambda} \ln(U)$$

where $U \sim \text{Unif}(0, 1)$

for a random variable Y with a pdf f_Y , if

$$f_{Y_1}(y_1) = e^{-y_1} \mathbf{1}_{y_1 \geq 0} = \begin{cases} e^{-y_1} & y_1 \geq 0 \\ 0 & o.w. \end{cases} = e^{-y_1} \mathbf{1}_{y_1 \geq 0} \Rightarrow Y_1 \sim \text{Exp}(-1)$$

Now consider the following change of variable, $Y_2 = -Y_1$, where the inverse transform is given by, $y_1 = -y_2$ and the jacobian is $\frac{dy_1}{dy_2} = -1$.

$$f_{Y_2}(y_2) = f_{Y_1}(-y_2) \cdot |-1| = \begin{cases} e^{y_2} & y_2 < 0 \\ 0 & o.w. \end{cases} = e^{y_2} \mathbf{1}_{y_2 < 0} \Rightarrow -Y_2 \sim \text{Exp}(-1)$$

$$\begin{aligned} g(x) &= 0.5e^x \mathbf{1}_{x < 0} + 0.5e^{-x} \mathbf{1}_{x \geq 0} \\ &= \frac{1}{2} \cdot f_{Y_2}(y_2) + \frac{1}{2} \cdot f_{Y_1}(y_1) \end{aligned}$$

where $Y_2 = -Y_1$ and $Y_1 \sim \text{Exp}(-1)$

Thus, we use the composition method by generating $Y_1 \sim \text{Exp}(-1)$ and multiply -1 with 0.5 probability. Y_1 can be generated with $Y_1 = \ln(U)$ where $U \sim \text{Unif}(0, 1)$.

```
size = int(1e5)

def sample_v(size:int) -> np.array:
    u1, u2 = np.random.random((2, size))
    v = np.log(u1)
    idx = u2 < 0.5
    v[idx] *= -1

    return v
```

```

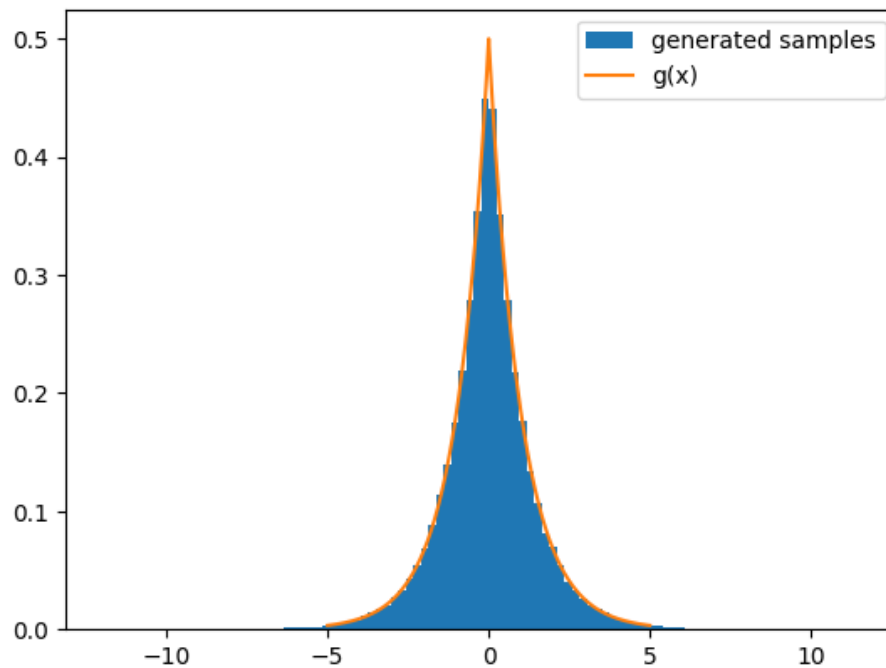
v = sample_v(size)

def g(x):
    return (x <= 0) * 0.5 * np.exp(x) + (x > 0) * 0.5 * np.exp(-x)

x = np.linspace(-5, 5, size)

plt.hist(v, bins=100, density=True, label='generated samples')
plt.plot(x, g(x), label="g(x)")
plt.legend()
plt.show()

```



```

a = np.sqrt(2 * np.e / np.pi)
def f(x):
    return np.exp(-(x ** 2) / 2) / 2 / np.sqrt(2 * np.pi)

```

```

def acc_rej(size):
    out = np.empty(size)
    i = num_unif = 0

    while i < size:
        v = sample_v(1)
        u = np.random.random()
        num_unif += 1

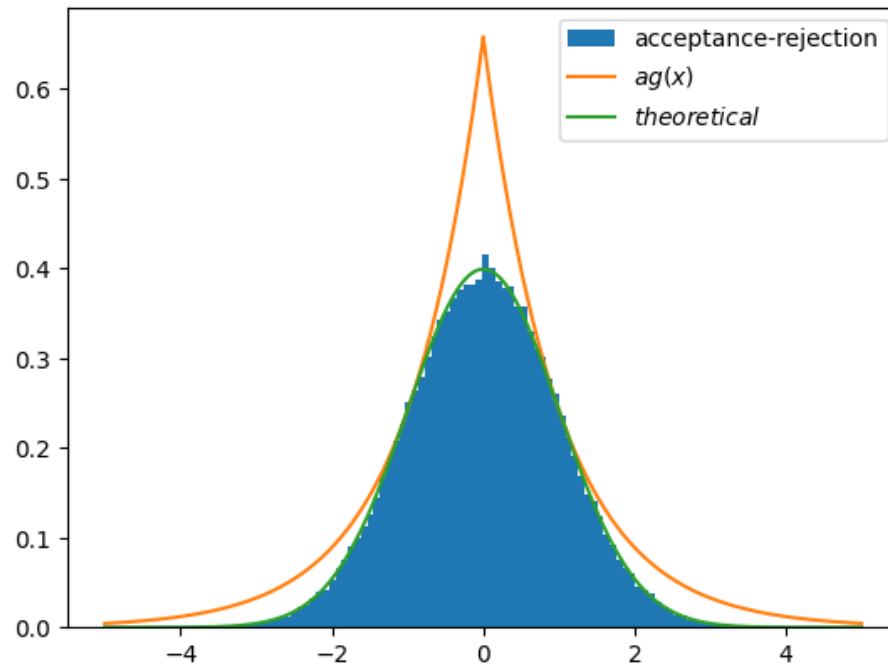
        if a * g(v) * u <= f(v):
            out[i] = v
            i += 1

    return out, num_unif

smpls, num_unif_ar = acc_rej(size)

plt.hist(smpls, bins=100, density=True, label='acceptance-rejection')
plt.plot(x, a * g(x), label=r'$ag(x)$')
plt.plot(x, f(x), label=r'$theoretical$')
plt.legend()
plt.show()

```



2 Box-Muller algorithm

```
def box_muller(n:int = 1000):
    out = np.empty(n)

    m = (n + 1) // 2
    size = (2, m)
    lamda = 1 / 2

    u1, u2 = np.random.random(size)
    num_unif = 2 * m

    v = -np.log(u1) / lamda
    w = u2 * 2 * np.pi

    x = np.sqrt(v) * np.cos(w)
    y = np.sqrt(v) * np.sin(w)

    out[:m] = x
    out[-m:] = y
```

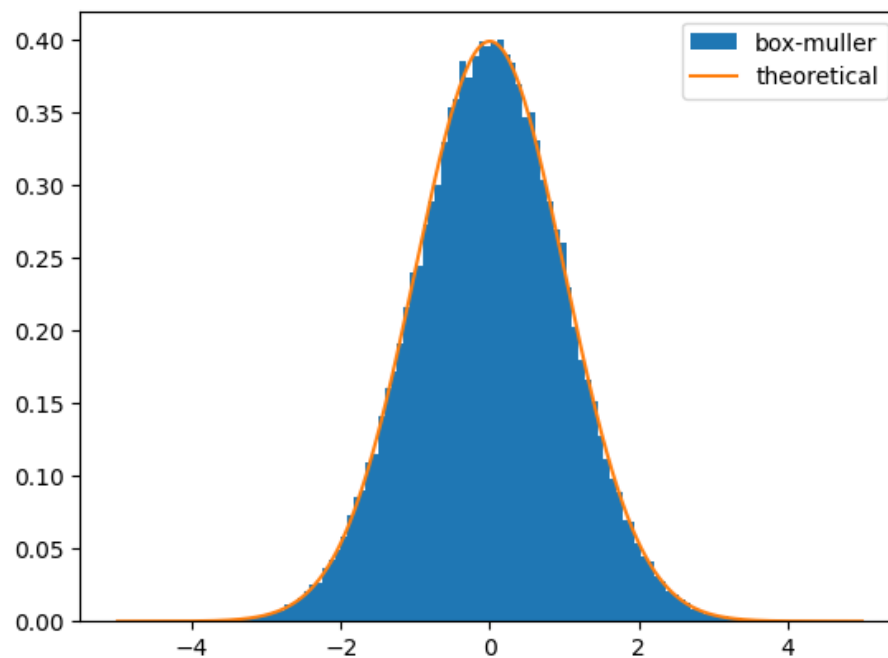
```

        return out, num_unif

smpls, num_unif_bm = box_muller(size)

plt.hist(smpls, 100, density=True, label = "box-muller")
plt.plot(x, f(x), label = "theoretical")
plt.legend()
plt.show()

```



3 Marsaglia-Bray

```

size = int(1e5)
def marsaglia_bray(size: int):
    out = np.empty(size)
    m = (size + 1) // 2

    tmp = np.empty((3, m))
    i = num_unif = 0

    while i < m:

```

```

    u1, u2 = np.random.uniform(-1, 1, size=2)
    num_unif += 2

    u = u1 ** 2 + u2 ** 2

    if u <= 1:
        tmp[:, i] = u1, u2, u
        i += 1

    u = tmp[2]
    z1, z2 = tmp[:2] * np.sqrt(-2 * np.log(u) / u)

    out[:m] = z1
    out[-m:] = z2

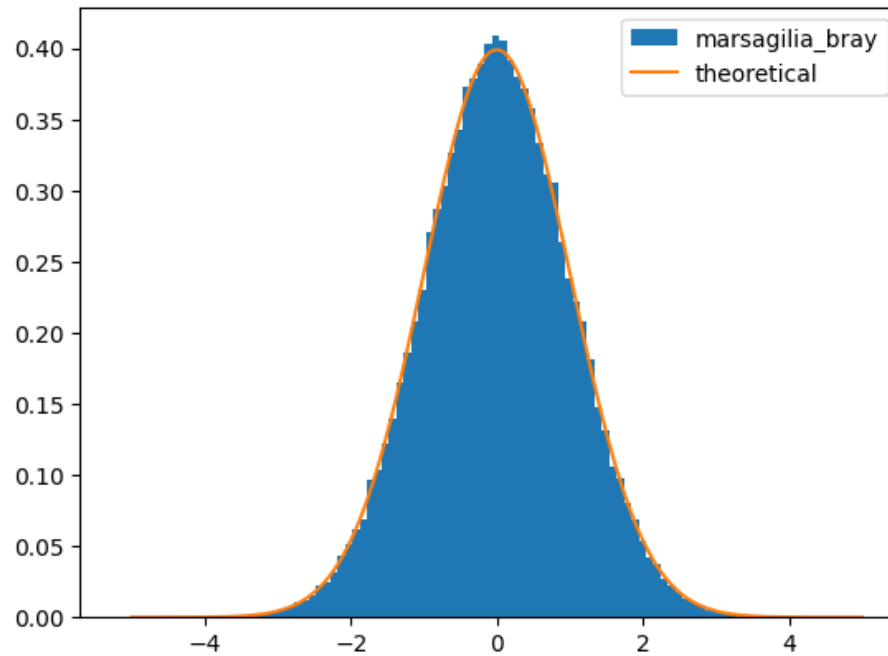
    return out, num_unif

smpls, num_unif_mb = marsagilia_bray(size)

x = np.linspace(-5, 5, size)

plt.hist(smpls, 100, density=True, label='marsagilia_bray')
plt.plot(x, f(x), label="theoretical")
plt.legend()
plt.show()

```



Acceptance-Rejection

```
result = %timeit -o acc_rej(size)
```

1.48 s \pm 6.73 ms per loop (mean \pm std. dev. of 7 runs, 1 loop each)

```
print(f"Avg number of uniform samples required : {num_unif_ar / size:.4f}")
print(f"Avg number of computation time required: {result.average / size}")
```

Avg number of uniform samples required : 1.3138

Avg number of computation time required: 1.4756457619999997e-05

Box-Muller

```
result = %timeit -o box_muller(size)
```

2.07 ms \pm 13.4 μ s per loop (mean \pm std. dev. of 7 runs, 100 loops each)

```
print(f"Avg number of uniform samples required : {num_unif_bm / size:.4f}")
print(f"Avg number of computation time required: {result.average / size}")
```


Avg number of uniform samples required : 1.0000
Avg number of computation time required: 2.0744293428571478e-08

Marsaglia-Bray

```
result = %timeit -o marsaglia_bray(size)
```

118 ms ± 907 µs per loop (mean ± std. dev. of 7 runs, 10 loops each)

```
print(f"Avg number of uniform samples required : {num_unif_mb / size:.4f}")  
print(f"Avg number of computation time required: {result.average / size}")
```

Avg number of uniform samples required : 1.2753
Avg number of computation time required: 1.1776846845714278e-06

Problem 2

Consider a three-state, discrete time Markov chain with states 1, 2 and 3 and transition matrix

$$P = \begin{pmatrix} 0.9 & 0.07 & 0.03 \\ 0.12 & 0.8 & 0.08 \\ 0.2 & 0.3 & 0.5 \end{pmatrix}$$

- 1) Using a uniformly distributed initial condition, generate one trajectory from $t = 0$ to $t = 100,000$ of this Markov chain and plot the first 500 time steps of the trajectory.
- 2) Use the full trajectory you calculated in part (1) to estimate the expected proportion of time spent in each state.
- 3) Calculate the dominant eigenvalue of P and its corresponding left eigenvector. How does this eigenvector relate to the expected times you found in part (2)?

```
np.random.seed(12)  
num_steps = int(1e5)  
  
p_mat = np.array([  
    [.9, .07, .03],  
    [.12, .8, .08],  
    [.2, .3, .5],  
)
```

```

traj = np.empty(num_steps, int)

unif_cdf = np.array([1/3, 2/3, 1])
u = np.random.random()
traj[0] = int(sum(u > unif_cdf))
x = np.eye(3)[traj[0]]

for i in range(1, num_steps):
    cdf = (x.T @ p_mat).cumsum()
    u = np.random.random()
    traj[i] = int(sum(u > cdf))
    x = np.eye(3)[traj[i]]

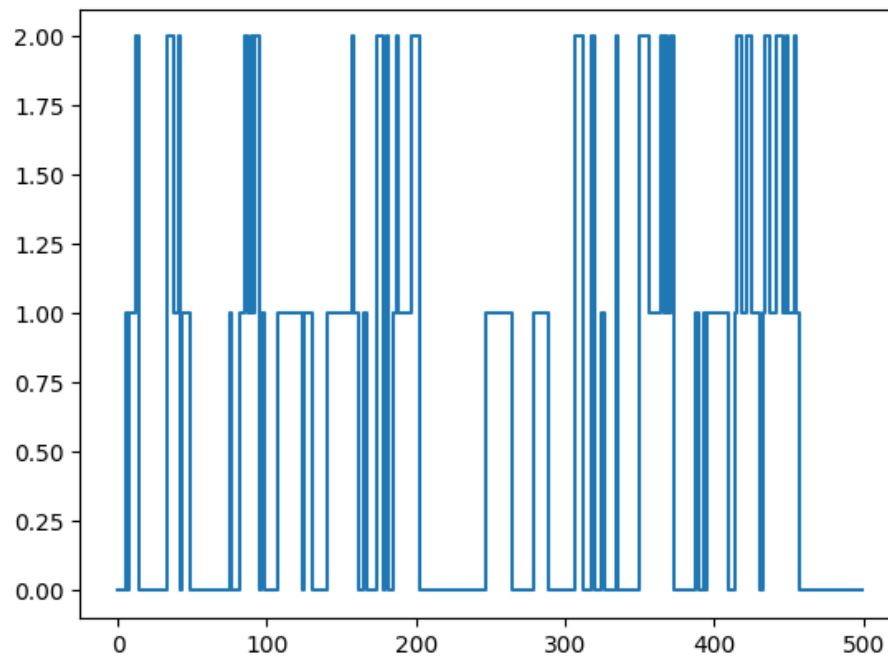
```

1

```

plt.step(np.arange(500), traj[:500], where='post')
plt.show()

```



2

```
_, cnt = np.unique(traj, return_counts=True)
freq = cnt / num_steps
freq.round(4)
```

array([0.5805, 0.3318, 0.0877])

3

```
eig_val, eig_vec = np.linalg.eig(p_mat.T)
stationary_dist = eig_vec[:, 0] / sum(eig_vec[:, 0])
print(f"Principal eigen value      : {eig_val[0]:.4f}")
print(f"corresponding eigen vector: {eig_vec[:, 0].round(4)}")
print(f"scaled eigen vector        : {stationary_dist.round(4)}")
```

```
Principal eigen value      : 1.0000
corresponding eigen vector: [0.858  0.4967 0.131 ]
scaled eigen vector        : [0.5775 0.3343 0.0881]
```

If we scale the principal eigenvector so that it sum up to 1, the scaled eigenvector is the stationary distribution of the markov chain. The long-term visiting frequency of each state or equivalently, the expected proportion of time spent in each state should converge to the stationary distribution as $t \rightarrow \infty$.

By definition of eigenvalue and eigenvector, we have

$$\mathbf{v}^T P = \lambda \mathbf{v}^T$$

where \mathbf{v} is the eigenvector with the corresponding eigenvalue of $\lambda = 1$

We can divide the BHS by the sum of all elements in a vector $a := \mathbf{1}^T \mathbf{v} = \sum v_i$. That is,

$$\begin{aligned} \mathbf{v}^T P &= \mathbf{v}^T \\ \frac{1}{a} \mathbf{v}^T P &= \frac{1}{a} \mathbf{v}^T \\ \pi^T P &= \pi^T \quad \text{where } \pi := \frac{1}{a} \mathbf{v} \end{aligned}$$

Since $\pi^T P = \pi^T P^n = \pi^T$, for $\forall n \in \{1, 2, \dots\}$, π is the stationary distribution of the Markov chain.

Problem 3

Suppose that an insurance company pays out $N \sim \text{Poisson}(10)$ claims per week and the payout for each claim is $X_i \sim \text{Exp}(0.1)$. (All of the random variables

N and X_i are mutually independent.) The total amount that the insurance company needs to pay out each week is given by

$$X = \sum_{i=1}^N X_i$$

Use the composition method to generate at least 100,000 i.i.d. samples of X . Plot a histogram of your samples and report the sample mean and variance.

```
from collections import defaultdict

def poisson(lamda:int , size: int) -> np.array:
    out = np.empty(size, int)
    cdf = defaultdict(float) # hashmap for a poisson cdf
    cdf[0] = np.exp(-lamda)

    for i in range(size):
        n = 0
        p = cdf[0]

        u = np.random.random()

        while u > cdf[n]:
            p = lamda * p / (n + 1)
            if cdf[n+1] == 0:
                cdf[n+1] += cdf[n] + p
            n += 1

        out[i] = n

    return out

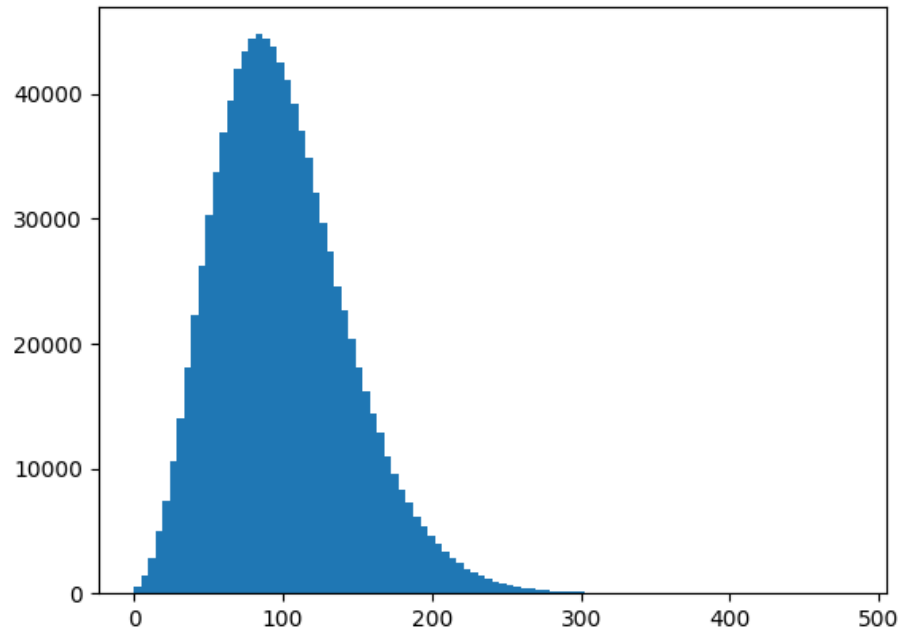
size = int(1e6)
n = poisson(10, size)

def exponential(lamda, size):
    u = np.random.random(size)
    return -np.log(u)/lamda

x = np.empty(size)
for i in range(size):
    x[i] = exponential(0.1, n[i]).sum()
```

```
print(f"Mean: {x.mean():.4f} \t Var: {x.var():.4f}")
plt.hist(x, 100)
plt.show()
```

Mean: 99.9678 Var: 1999.9063



Problem 4

Consider the multivariate normal random variable $X \sim N(\mathbf{0}, \Sigma)$, where

$$\Sigma = \begin{pmatrix} 1 & -1 \\ -1 & 1 \end{pmatrix}$$

- 1) What is the Cholesky decomposition of Σ ? Explain why it might be problematic to generate samples of X using this decomposition.
- 2) What is the pdf of X ?
- 3) Generate 1,000 samples of X and make a scatter plot of the points. How does this plot relate to your answer from parts (1) and (2)?

1

The characteristic polynomial of Σ is given by,

$$\begin{aligned} |\Sigma - \lambda I| &= \begin{vmatrix} 1 - \lambda & -1 \\ -1 & 1 - \lambda \end{vmatrix} \\ &= \det \begin{pmatrix} 1 - \lambda & -1 \\ -1 & 1 - \lambda \end{pmatrix} \\ &= (1 - \lambda)^2 - 1 \end{aligned}$$

The eigenvalue of Σ is the solutions of the characteristic equation

$$(1 - \lambda)^2 - 1 = 0$$

$$\lambda_1 = 2, \lambda_2 = 0$$

Since λ_2 is 0, Σ is not positive definite, but only positive semidefinite.

Thus, we can get the following LDL^T decomposition.

$$\begin{aligned} \Sigma &= LDL^T \\ &= \begin{pmatrix} 1 & 0 \\ L_{21} & 1 \end{pmatrix} \begin{pmatrix} D_1 & 0 \\ 0 & D_2 \end{pmatrix} \begin{pmatrix} 1 & L_{12} \\ 0 & 1 \end{pmatrix} \\ &= \begin{pmatrix} D_1 & D_1 L_{12} \\ L_{21} D_1 & L_{21} D_1 L_{12} + D_2 \end{pmatrix} \\ \begin{pmatrix} D_1 & D_1 L_{12} \\ L_{21} D_1 & L_{21} D_1 L_{12} + D_2 \end{pmatrix} &= \begin{pmatrix} 1 & -1 \\ -1 & 1 \end{pmatrix} \end{aligned}$$

$$D_1 = 1, D_2 = 0, L_{12} = L_{21} = -1$$

That is,

$$\begin{aligned} \Sigma &= \begin{pmatrix} 1 & 0 \\ -1 & 1 \end{pmatrix} \begin{pmatrix} 1 & 0 \\ 0 & 0 \end{pmatrix} \begin{pmatrix} 1 & -1 \\ 0 & 1 \end{pmatrix} \\ &= \begin{pmatrix} 1 & 0 \\ -1 & 1 \end{pmatrix} \begin{pmatrix} 1 & 0 \\ 0 & 0 \end{pmatrix} \begin{pmatrix} 1 & 0 \\ 0 & 0 \end{pmatrix} \begin{pmatrix} 1 & -1 \\ 0 & 1 \end{pmatrix} \\ &= \begin{pmatrix} 1 & 0 \\ -1 & 0 \end{pmatrix} \begin{pmatrix} 1 & -1 \\ 0 & 0 \end{pmatrix} \\ &= AA^T \end{aligned}$$

If we sample X from this decomposition A , That is, $X = \mu + AZ$ where $Z = (Z_1, Z_2)^T \sim N(0, I)$

$$X = \begin{pmatrix} 0 \\ 0 \end{pmatrix} + \begin{pmatrix} 1 & 0 \\ -1 & 0 \end{pmatrix} \begin{pmatrix} Z_1 \\ Z_2 \end{pmatrix} = \begin{pmatrix} Z_1 \\ -Z_1 \end{pmatrix}$$

Thus, this method discards all the generated z_2 and this will cause a significant computational inefficiency. Because this method is generating the twice of the number of samples that is acutally needed.

2

X does not have a two dimensional density, but only has a one dimensional normal density as follows.

$$f_X(\mathbf{x}) = \begin{cases} \frac{1}{\sqrt{2\pi}} \exp\left(-\frac{x_1^2}{2}\right) & \text{if } x_2 = -x_1 \\ 0 & \text{o.w} \end{cases}$$

3

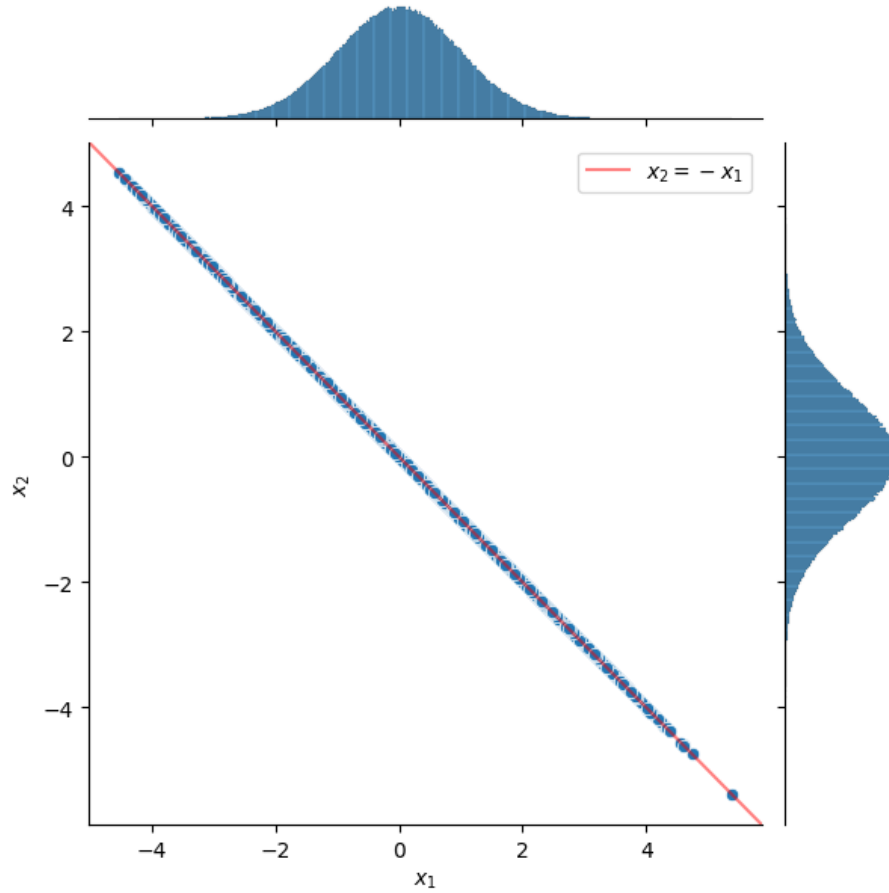
```
a_mat = np.array([[1, 0], [-1, 0]])

x = np.empty((size, 2))

for i in range(size):
    z, _ = box_muller(2)
    x[i] = a_mat @ z.T

import seaborn as sns

sns.jointplot(x=x[:, 0], y=x[:, 1])
plt.axline([0,0], [1, -1], color="red", alpha=.5, label=r"$x_2 = -x_1$")
plt.xlabel(r"$x_1$")
plt.ylabel(r"$x_2$")
plt.legend()
plt.show()
```



The generated samples lie on the negative identity line, indicating that the support of the joint distribution of X is a line of $X_2 = -X_1$. That is, random vector X has a one dimensional standard normal density, which is consistent with the result of the part 1 and 2.

Problem 5

Suppose $S \sim GBM(0.1, 0.09)$. In this problem, you will simulate S_{t_k} at the times $t_k = 0.01k$ for $0 \leq k \leq 1000$. Use the initial condition $S_0 = 1$.

- 1) Simulate S once and plot the trajectory.
- 2) Simulate S at least 10,000 times and plot the mean trajectory.
- 3) Estimate the probability that the maximum of S_{t_k} is no greater than 6.
- 4) Estimate the probability that S_{t_k} ever leaves the interval $[0.25, 10]$.


```

np.random.seed(12)

mu = 0.1
sigma = 0.09

dt = 0.01
T = 1000
t = np.arange(T+1) * dt
S_0 = 1

size = 10000

s = np.ones((size, len(t)))
s[:, 0] = S_0
s_T = np.empty(size)

mean = (mu - sigma ** 2 / 2) * dt
sd = sigma * np.sqrt(dt)

for k in range(size):
    z, _ = box_muller(1000)
    log_diff = mean + sd * z
    s[k, 1:] = s[k, 0] * np.exp(log_diff).cumprod()
    s_T[k] = s[k, -1]

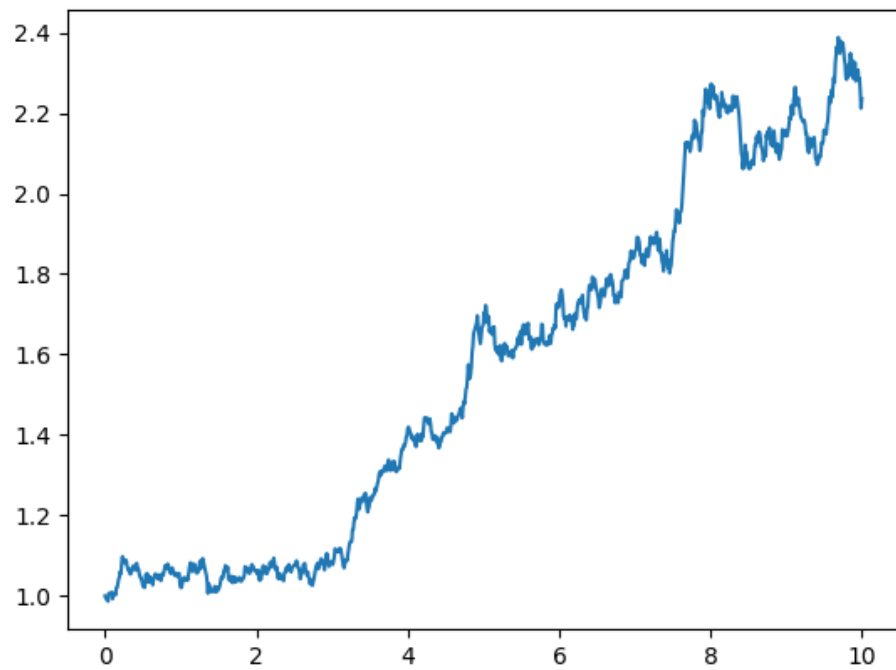
```

1

```

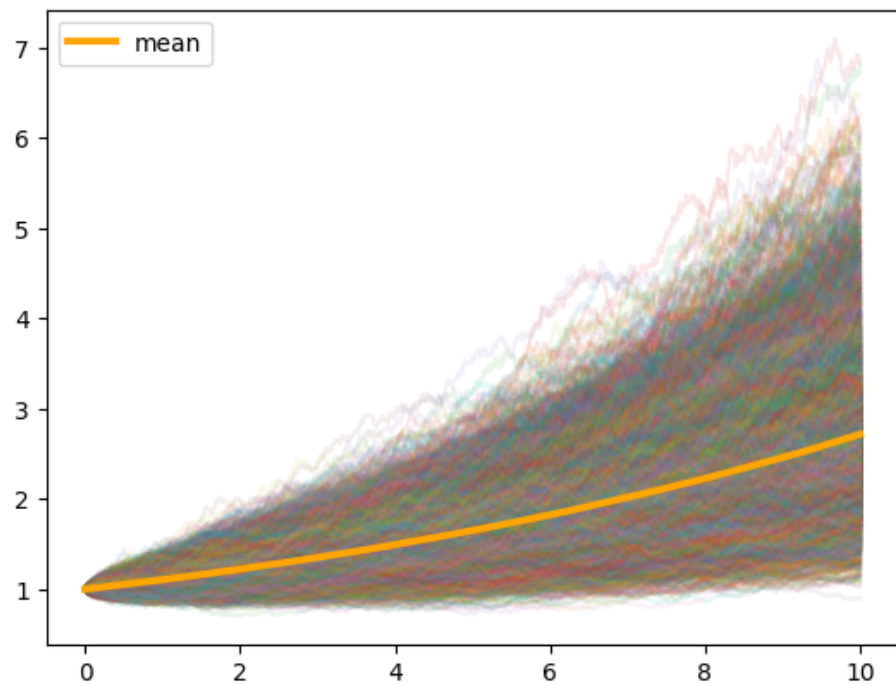
plt.plot(t, s[0])
plt.show()

```



2

```
for i in range(size):  
    plt.plot(t, s[i], alpha=0.1)  
plt.plot(t, s.mean(0), color='orange', lw=3, label='mean')  
plt.legend()  
plt.show()
```



3

$$P[\max(S_{t_k}) \leq 6] =$$

```
np.mean(s.max(1) <= 6)
```

0.9984

4

$$P[\min(S_{t_k}) < 0.25 \text{ or } \max(S_{t_k}) > 10] = 1 - P[\min(S_{t_k}) \geq 0.25 \text{ and } \max(S_{t_k}) \leq 10] =$$

```
1 - np.mean((s.min(1) >= 0.25) * (s.max(1) <= 10))
```

0.0